

Machine learning

Feature Selection

Lecture V

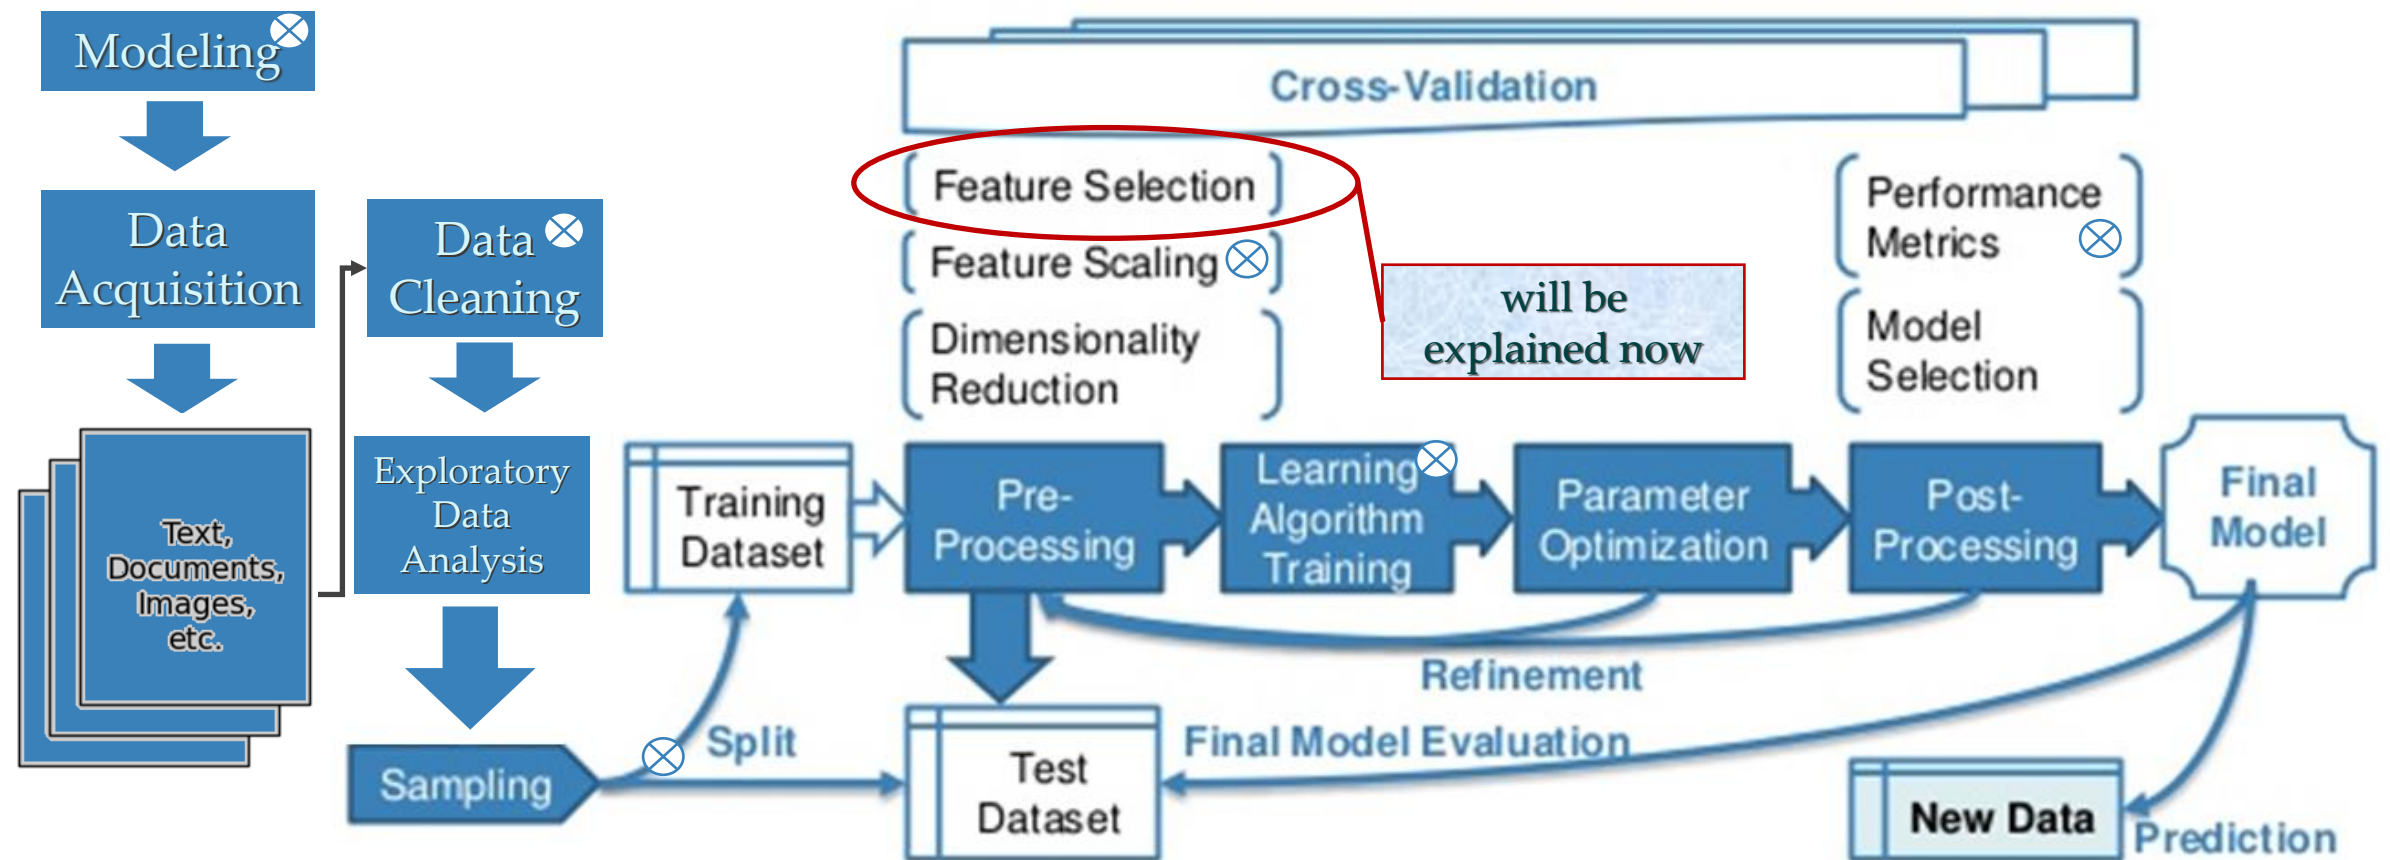
פיתוח:
ד"ר יהונתן שלר
משה פרידמן

What will we talk about

- ❖ A typical classification flow summary
- ❖ Feature selection

A typical classification flow

- diving in



Data Cleaning

- Duplicates
- Missing Data
- Remove
- Repair

Train-Test split

- + Validation-set

Data Exploration

Scaling

- Minmax
- norm.
- t-dist.
- standardization

Learning Algos.

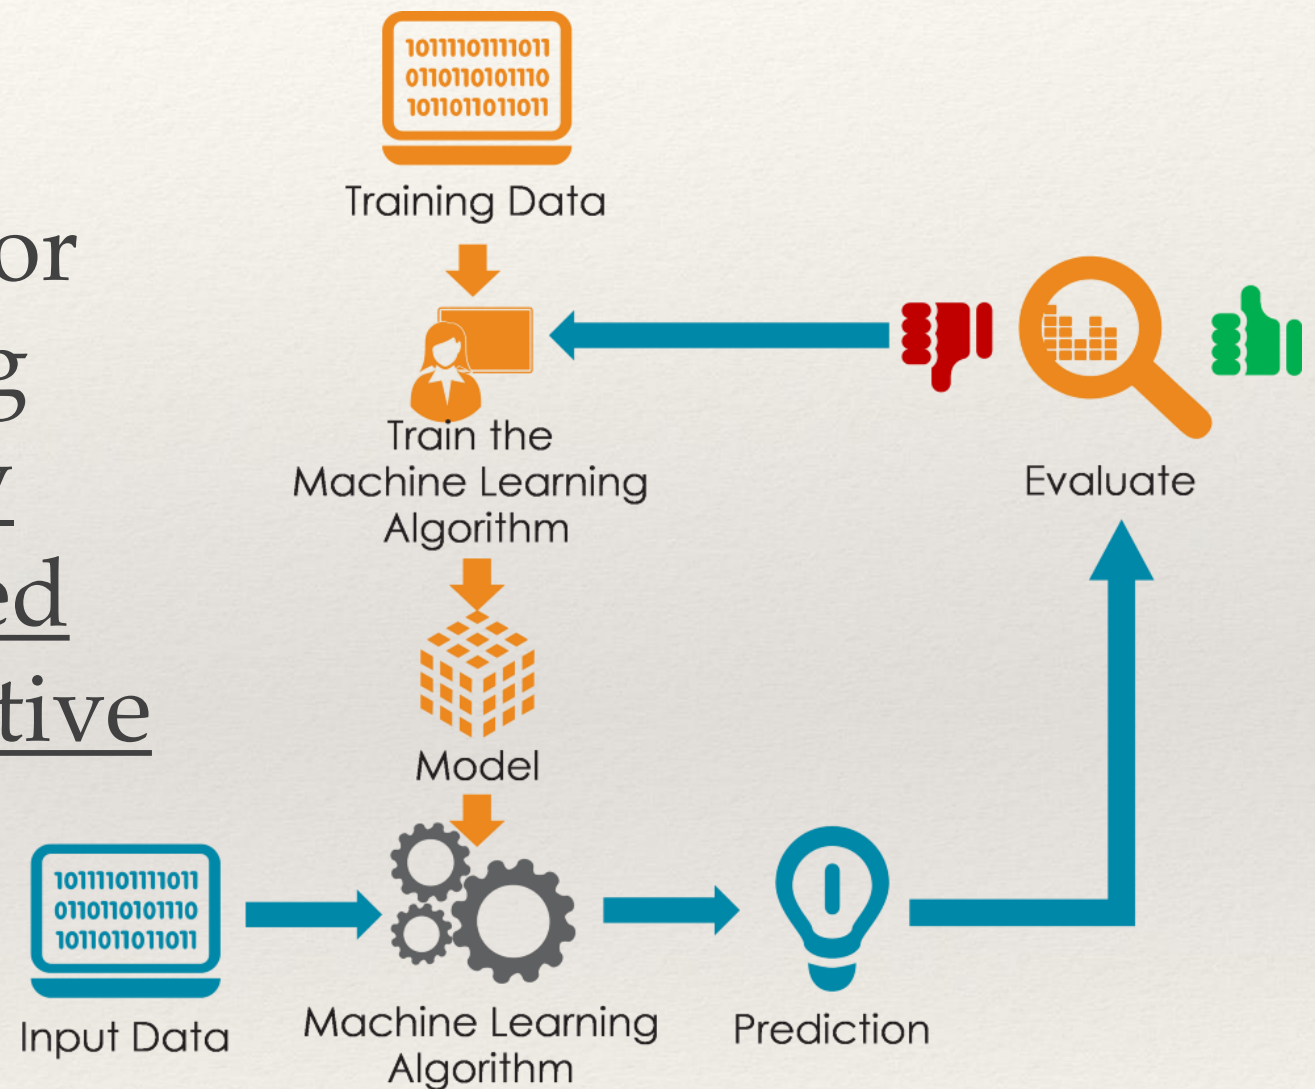
- KNN
- Decision Trees
- Naïve Bayes

Evaluation

- Confusion matrix
- Accuracy ,Error (rate)
- Precision, Recall
- F1 (soon)

Machine learning training

- ❖ A machine learning algorithm (e.g., classification, regression or clustering) uses a training dataset to determine how can the features be applied to unseen data for predictive purposes.



What is feature selection

- ❖ feature selection is the process of selecting a subset of relevant features for use in model construction” or in other words, the selection of the most important features

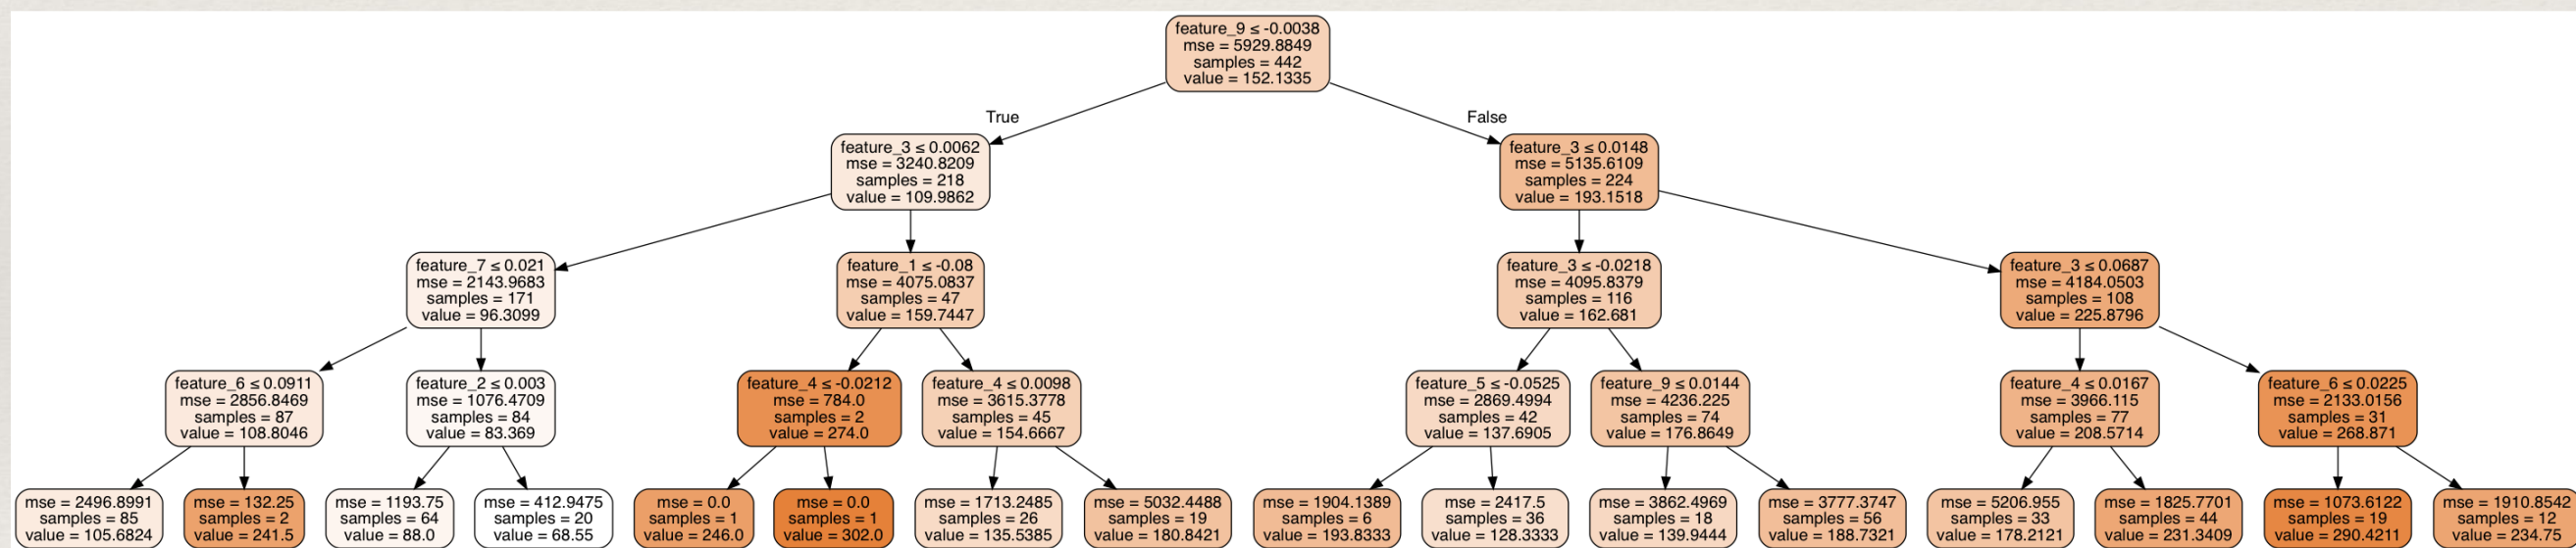
All features



Selected Features

Feature selection and Dimensionality Reduction - Motivation

Prevents Overfitting: A high-dimensional dataset having too many features can sometimes lead to overfitting (model captures both real and random effects).



For example – a decision tree could over-fit a trainset with too many features

Feature selection and Dimensionality Reduction - Motivation

Prevents Overfitting: A high-dimensional dataset having too many features can sometimes lead to overfitting (model captures both real and random effects).

Simplicity: An over-complex model having too many features can be hard to interpret.

Feature selection & Dimensionality Reduction - Motivation

Prevents Overfitting: A high-dimensional dataset having too many features can sometimes lead to overfitting (model captures both real and random effects).

Simplicity: An over-complex model having too many features can be hard to interpret especially when features are correlated with each other.

Computational Efficiency: A model trained on a lower-dimensional dataset is computationally efficient (execution of algorithm requires less computational time).

Feature selection – techniques

1. not complex enough for learning

1. a. A trivial case – constant value (variance = 0)

→ We mentioned this example

1. b. Remove features with low variance

❖ Features with very low variance (under a threshold) are not complex enough for learning.

Feature selection – techniques

2. highly correlated features

2. Remove highly correlated features

- ❖ High correlated features could cause distortion of distance functions (KNN).
- ❖ Features that are highly correlated or co-linear can cause overfitting (NB)
- ❖ When a pair of variables are highly correlated, we can remove one without much loss of information.
 - ❖ Which one should we keep?
 - ❖ The one with a higher correlation to the target (see ahead)

Feature selection – techniques

2. highly correlated features

2. Remove highly correlated features

- ❖ Features that are highly correlated or co-linear can cause overfitting.
- ❖ zero imply weak or no correlation
- ❖ coefficients are used to measure the strength of the relationship between two variables.
- ❖ A trivial case – duplicate features

Feature selection – techniques

2.a. highly correlated features - Pearson correlation

2. Remove highly correlated features

- ❖ Features that are highly correlated or co-linear can cause overfitting.
- ❖ Pearson correlation is the one most commonly used
- ❖ Linear correlation
- ❖ Values always range:-1 (strong negative relationship) and +1 (strong positive relationship). 0 – no correlation.

Feature selection – techniques

2.a. highly correlated features - Pearson correlation

Covariance in the population:

$$Cov(X, Y) = \frac{\sum_{i=1}^n (x_i - \bar{X}) \cdot (y_i - \bar{Y})}{n}$$

Where:

X_i – the values of the X-variable

Y_j – the values of the Y-variable

\bar{X} – the mean (average) of the X-variable

\bar{Y} – the mean (average) of the Y-variable. הקלד משוואה כאן.

n – the number of data points

Covariance in the sample:

$$Cov(X, Y) = \frac{\sum_{i=1}^n (x_i - \bar{X}) \cdot (y_i - \bar{Y})}{n-1}$$

Positive covariance: Indicates that two variables tend to move in the same direction.

Negative covariance: Reveals that two variables tend to move in inverse directions.

Covariance Measures relationship, not strength

Feature selection – techniques

2. a. highly correlated features - Pearson correlation (2)

2. Remove highly correlated features

- ❖ Features that are highly correlated or co-linear can cause overfitting.

Pearson correlation - between x,y

$$\rho_{X,Y} = \frac{\text{cov}(X, Y)}{\sigma_X \sigma_Y}$$
$$= \frac{\sum_i (x_i - \bar{x})(y_i - \bar{y})}{\sqrt{\sum_i (x_i - \bar{x})^2 \sum_i (y_i - \bar{y})^2}}$$

Feature selection – techniques

2. highly correlated features – NMI

2. Remove highly correlated features

- ❖ Mutual information between two features (f_1, f_2)

Normalized Mutual information

$$\text{NMI} = \frac{IG(f_1; f_2)}{|H(f_1) + H(f_2)|/2}$$


- ❖ NMI value close to 1
→ high similarity
- ❖ NMI value close to 0
→ high dissimilarity

Feature selection – techniques

3. features with high correlation to target

3. Select features with high correlation to target

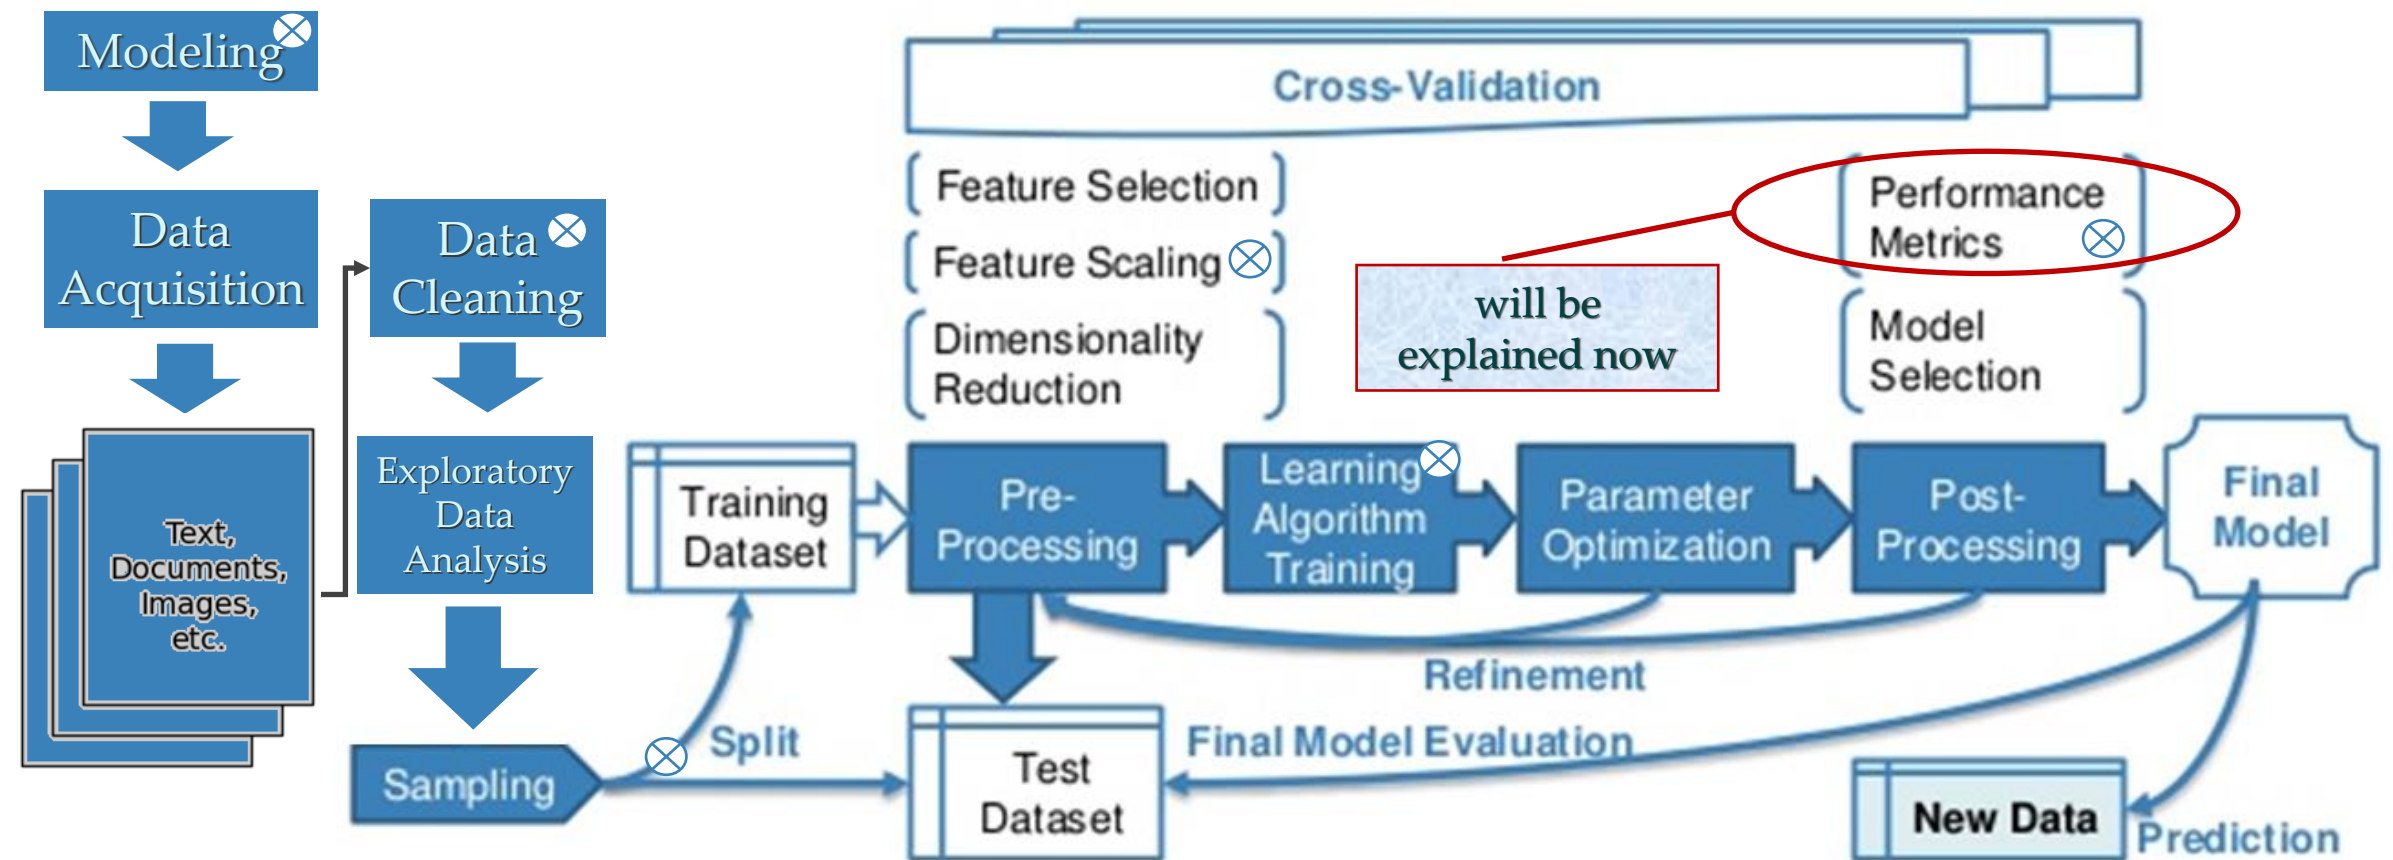
- ❖ Features that are high correlation to the class.
- ❖ How to use?
 - ❖ Choose top k features
 - ❖ Choose features passing threshold
- ❖ E.g., Mutual information based



$$NMI = \frac{IG(X|Y)}{|H(X) + H(Y)|/2}$$

A typical classification flow

- diving in



Data Exploration

Data Cleaning

- Duplicates
- Missing Data
- Remove
- Repair

Train-Test split

+ Validation-set

Scaling

- Minmax
- norm.
- t-dist.
- standardization

Learning Algos.

- KNN
- Decision Trees
- Naïve Bayes

Evaluation

- Confusion matrix
- Accuracy ,Error (rate)
- Precision, Recall
- F1 (soon)

Classification Measures – Confusion Matrix - Reminder

	Class 1 Predicted	Class 2 Predicted
Class 1 Actual	TP	FN
Class 2 Actual	FP	TN

- Class 1: Positive
- Class 2: Negative
- TP: True Positive
- FN: False Negative
- FP: False Positive
- TN: True Negative

- Classification Rate / Accuracy:

$$\frac{TP + TN}{TP + TN + FP + FN}$$

- Precision

$$\frac{TP}{TP + FP}$$

- High recall, low precision: Most of the positive examples are correctly recognized (low FN) but there are a lot of false positives.

- Recall

$$\frac{TP}{TP + FN}$$

- Low recall, high precision: We miss a lot of positive examples (high FN) but those we predict as positive are indeed positive (low FP).

Classification Measures – F1 Score

- It is useful to have one number to measure the performance of the classifier
- $F_{\alpha} = (1 + \alpha^2) \frac{Precision * Recall}{\alpha^2 * Precision + recall}$
- When $\alpha=1 \rightarrow F_1 = 2 \frac{Precision * Recall}{Precision + recall}$


התפלגות נורמלית

התפלגות נורמלית: נקראת גם גאוסיאן (Gaussian) או עקומת פעמון.


❖ פונקציית צפיפות סמטרית.

התפלגות z : תת קבוצה של התפלגות נורמלית בו התוחלת/הממוצע $=0$ וסטיית התקן $=1$.

❖ כל התפלגות נורמלית ניתן להפוך להתפלגות z

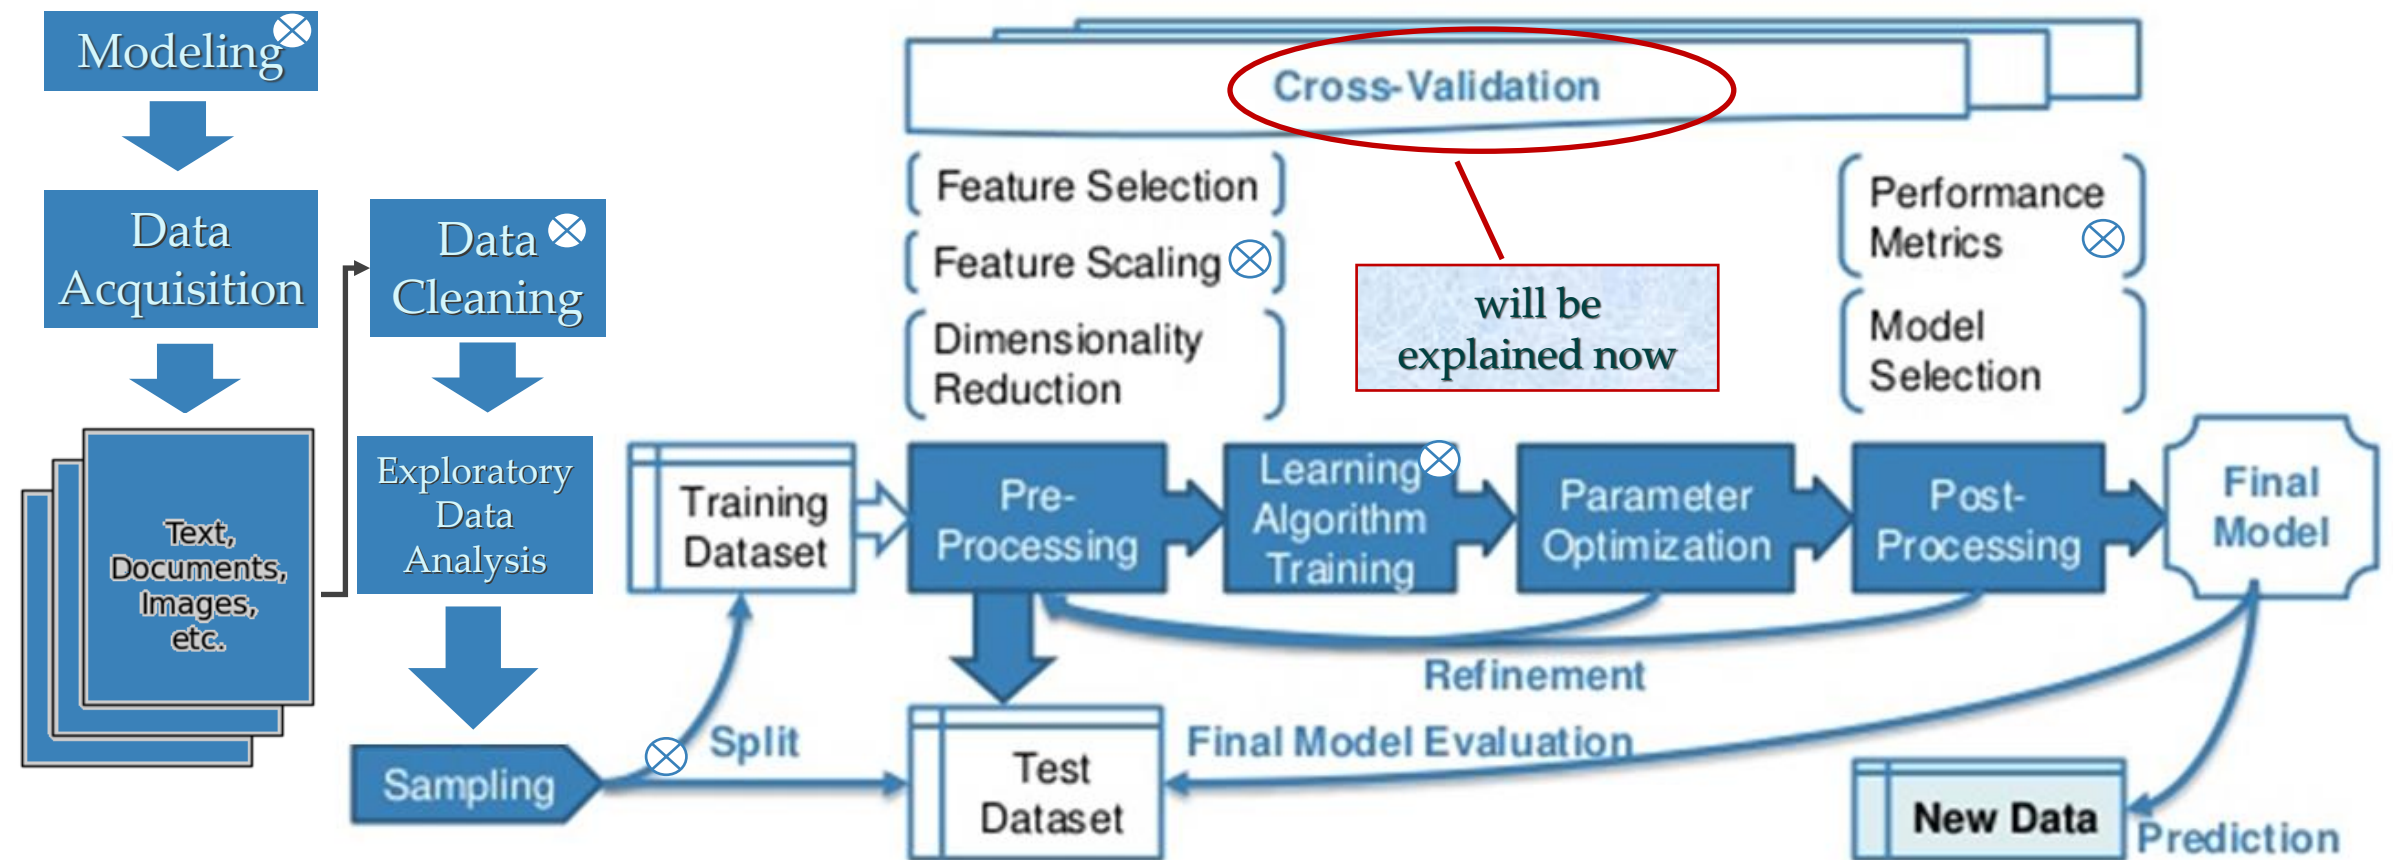


Outlier detection



A typical classification flow

- diving in



Data Exploration

Data Cleaning

- Duplicates
- Missing Data
- Remove
- Repair

Train-Test split

+ Validation-set

Scaling

- Minmax norm.
- t-dist. standardization

Learning Algos.

- KNN
- Decision Trees
- Naïve Bayes

Evaluation

- Confusion matrix
- Accuracy, Error (rate)
- Precision, Recall
- F1 (soon)


נושאים

Overfitting ❖

Model selection ❖


Validation ❖

Overfitting




Overfitting

- *Given a hypothesis space H , $h \in H$ overfits the training data if there exists some alternative hypothesis $h' \in H$ such that h has smaller error than h' over the training examples, but h' has smaller error than h over the entire distribution of instances.*



- Overfitting: Small error on training set, but large error on unseen examples.
- Underfitting: Larger error on training and test sets.

Overfitting

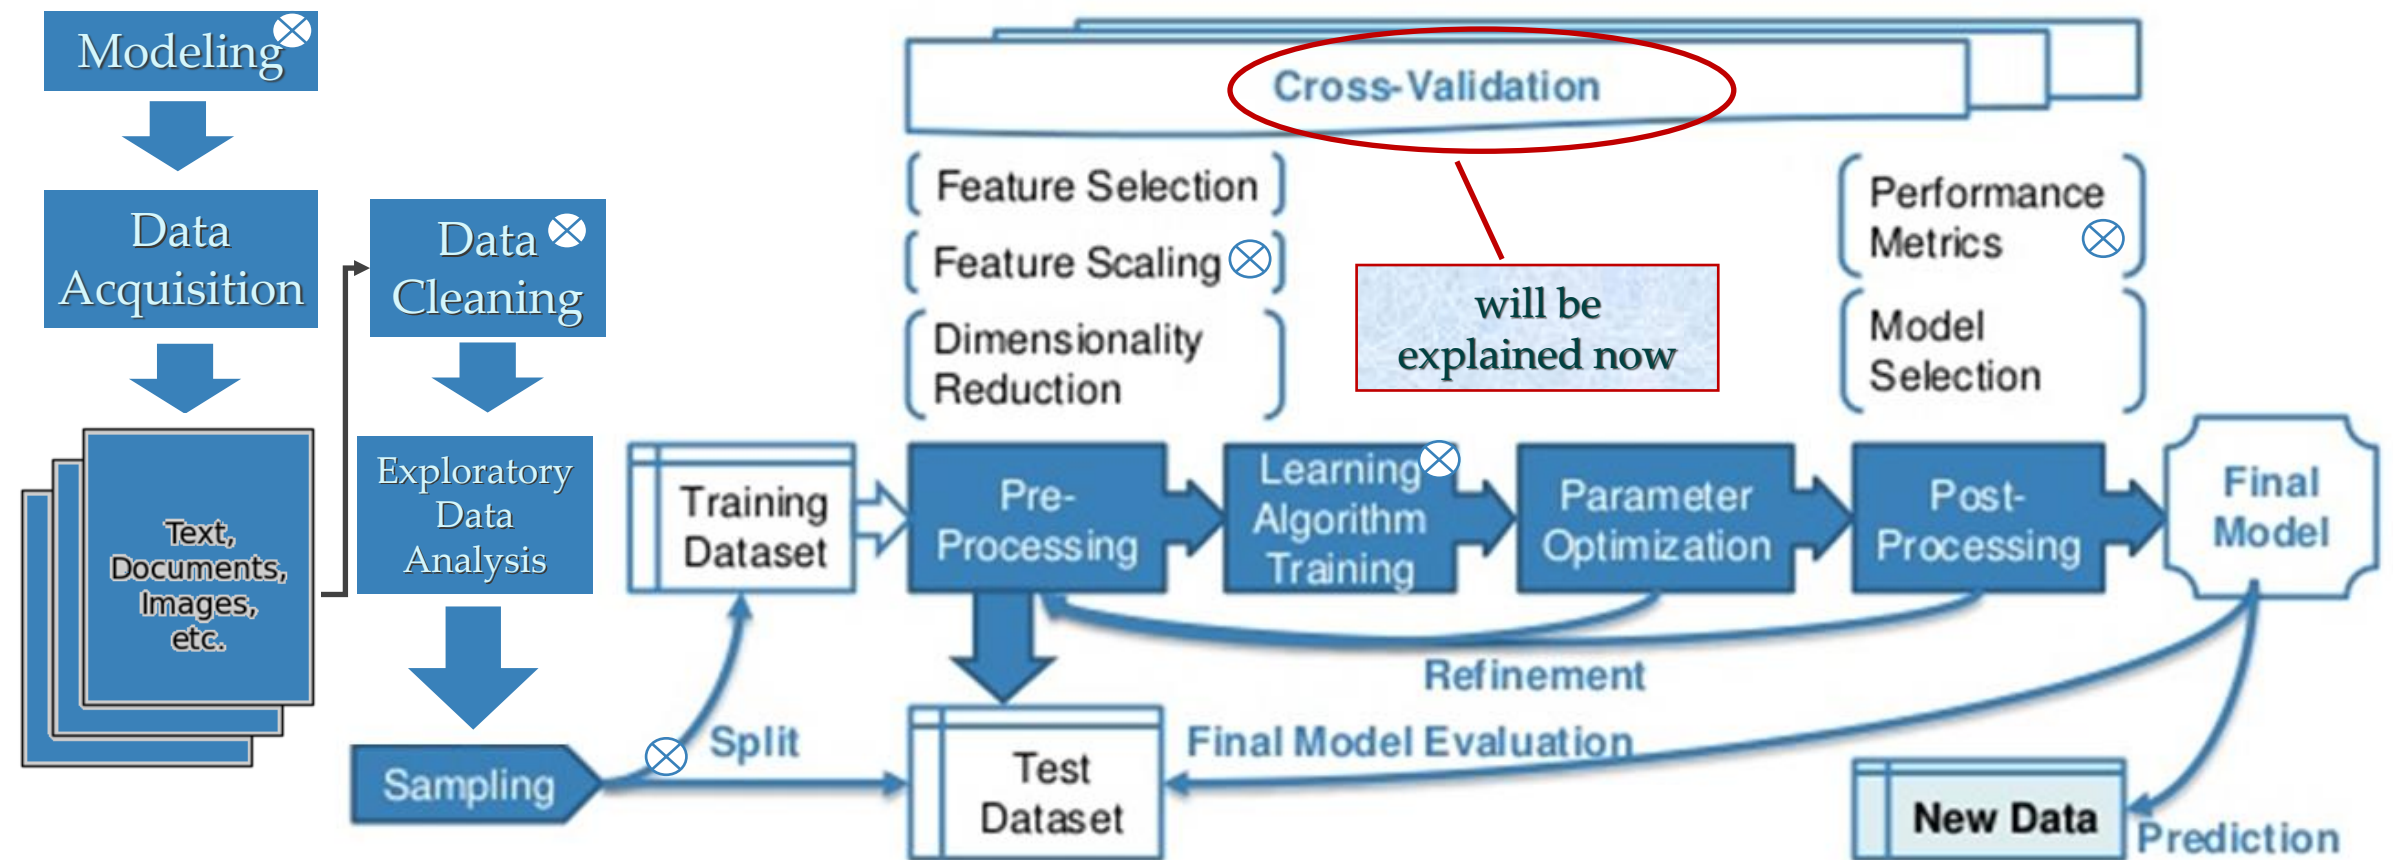


(by Tomaso Poggio, <http://www.mit.edu/~9.520/spring12/slides/class02/class02.pdf>)

- Green: True target function
- Red: Training points
- Blue: What we have learned (overfitting)

- The algorithm has learned perfectly the training examples, even the noise present in the examples and cannot generalise on unseen examples.

A typical classification flow



Data Cleaning

- Duplicates
- Missing Data
- Remove
- Repair

Train-Test split

+ Validation-set

Scaling

- Minmax norm.
- t-dist. standardization

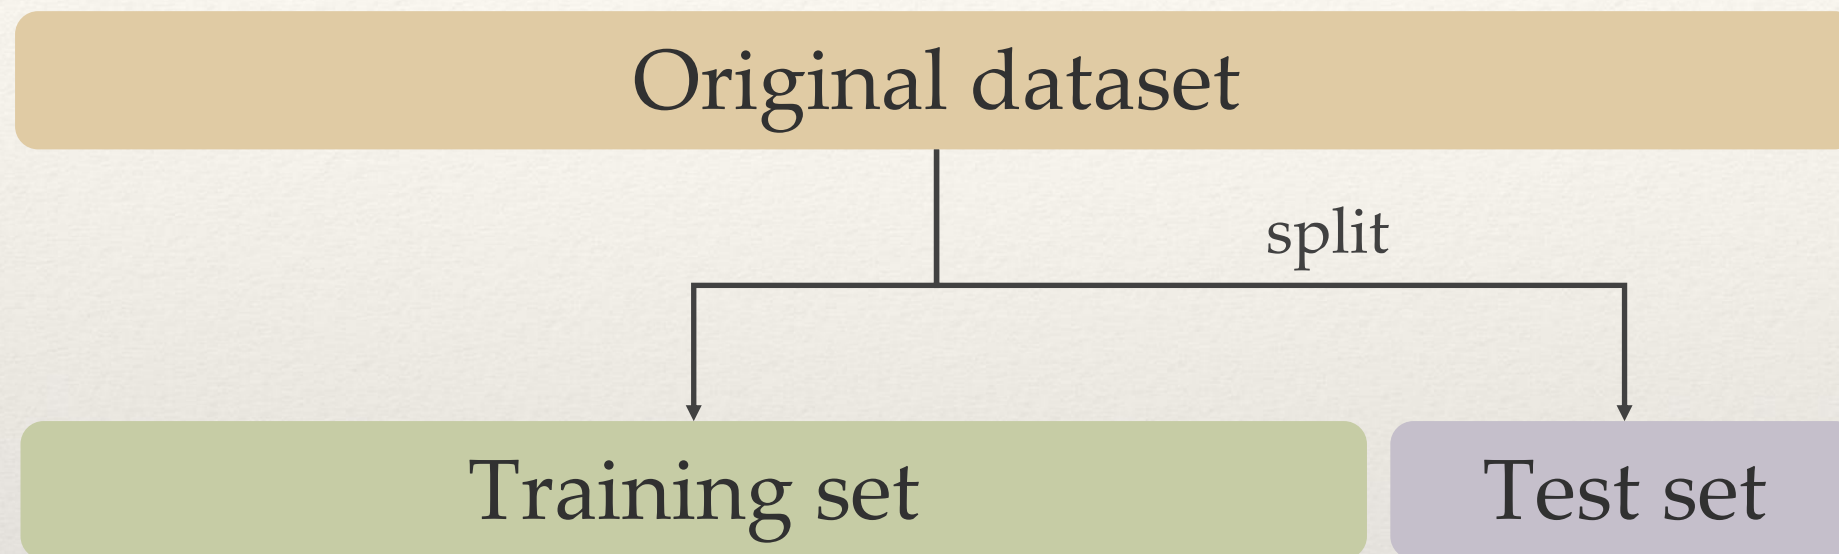
Learning Algos.

- KNN
- Decision Trees

Evaluation

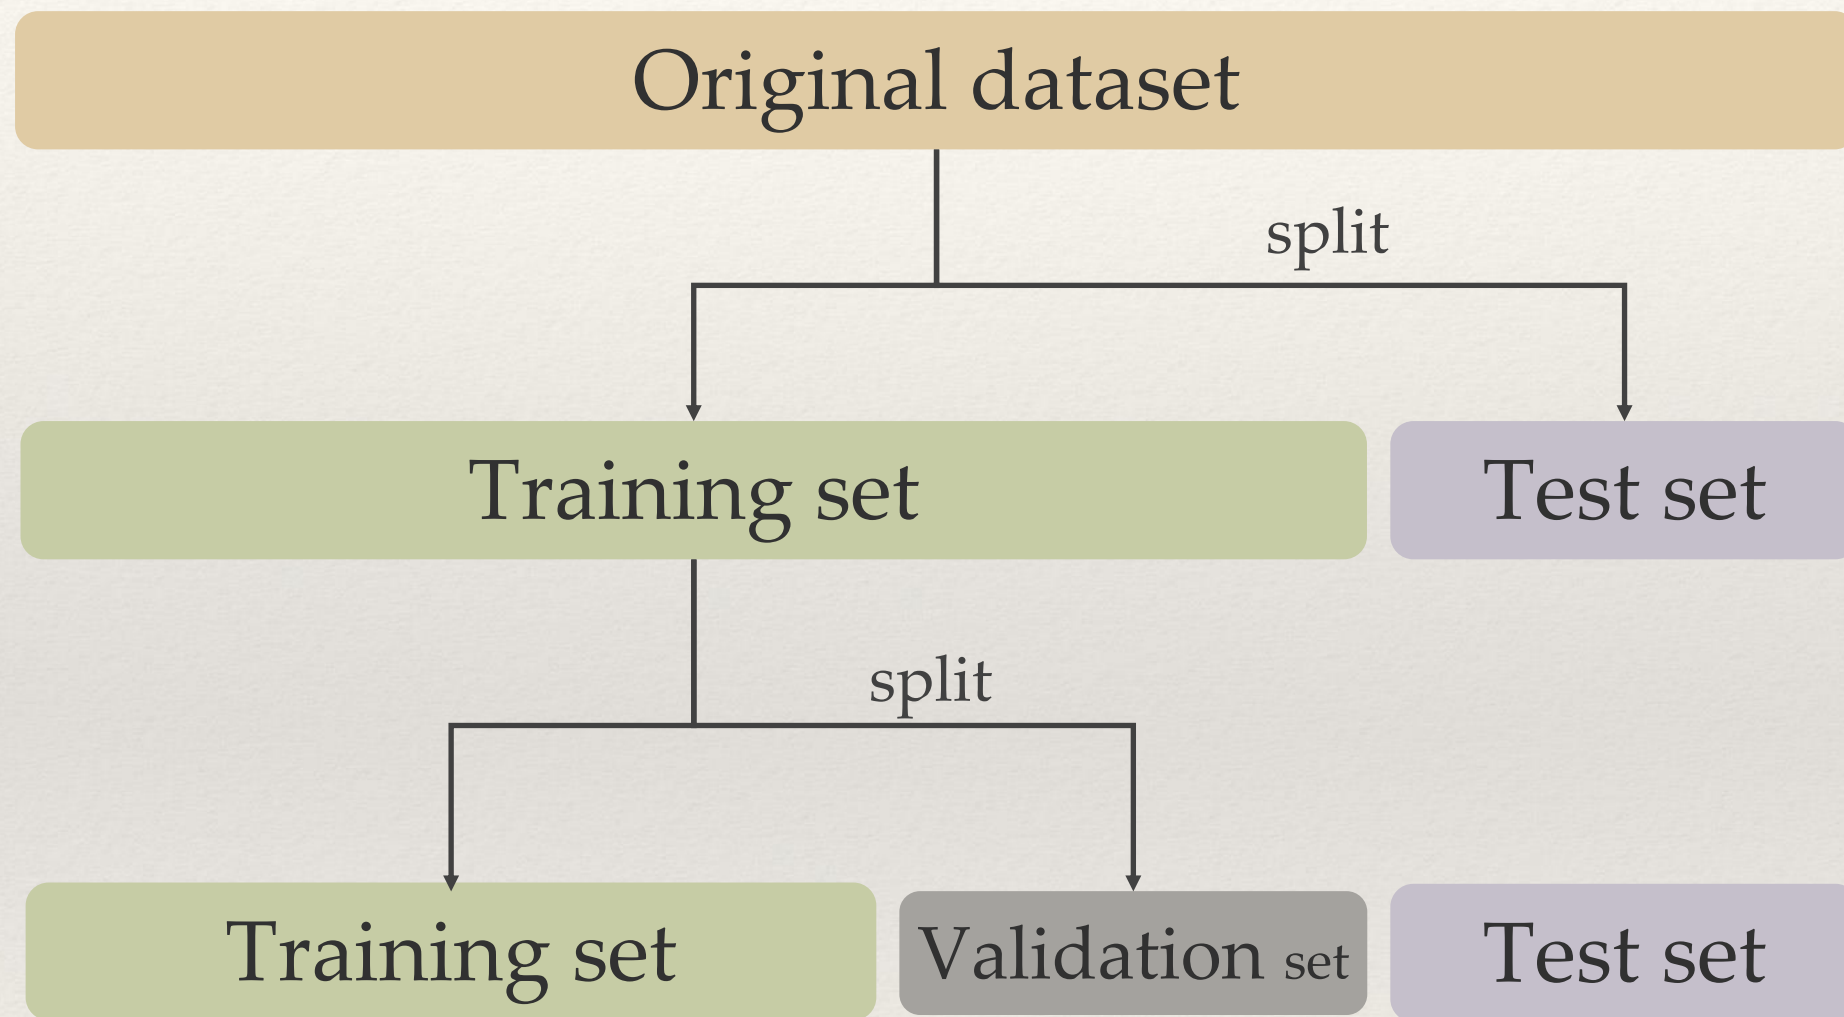
- Error (rate)
- Accuracy
- Confusion matrix

dataset – train-set and test-set



Validation set

dataset – train-test-validation



שימושים ל validation set

- ❖ סיוע במניעת overfitting
 - ❖ Model selection
 - ❖ בחירת hyperparameters מיטביים
 - ❖ תהליכים משלימים לתהליך האימון
 - ❖ Post pruning של עצי החלטה
 - ❖ שיערוך המודל, בהיעדר test מקובל cross validation (בהמשך ...)
- ועוד ...

What is Model Selection?

Given a set of models $M=\{M_1, M_2, \dots, M_R\}$, choose the model that is expected to do the best on the test data. The set M may consist of:

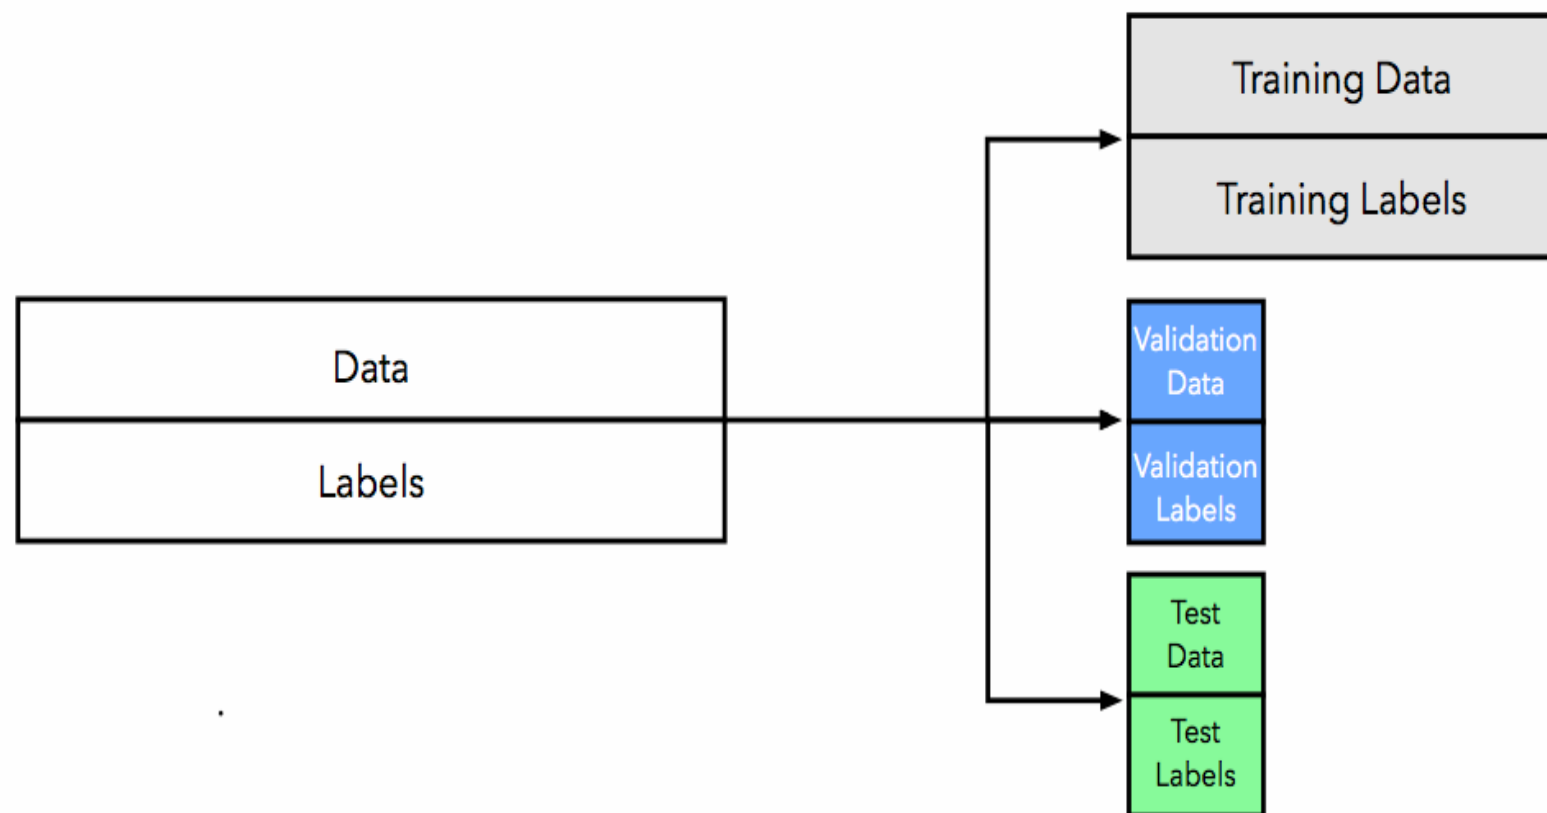
- Instances of same model with different complexities or hyperparams.
E.g.,
- K-Nearest Neighbors: Different choices of K
- Decision Trees: Different choices of the number of levels/leaves
- Architecture of a deep neural network (# of layers, nodes in each layer, activation function, etc)
- Naïve Bayes – smoothing methods
- Different types of learning models (e.g. KNN, DT, etc.)

Hyperparameters שלמדנו

- ❖ kNN - שיערוך k , שיטת מרחק, p בשיטת Minikowski distance, משקול מרחקי הנקודות
- ❖ עצי החלטה – עומק מקסימלי, מינימום דוגמאות בעלה, כמות המאפיינים לבדיקה כשמחפשים split מסוים
- ❖ Naive Bayes - פרמטרי החלקה, שיטות החלקה

Holdout Method

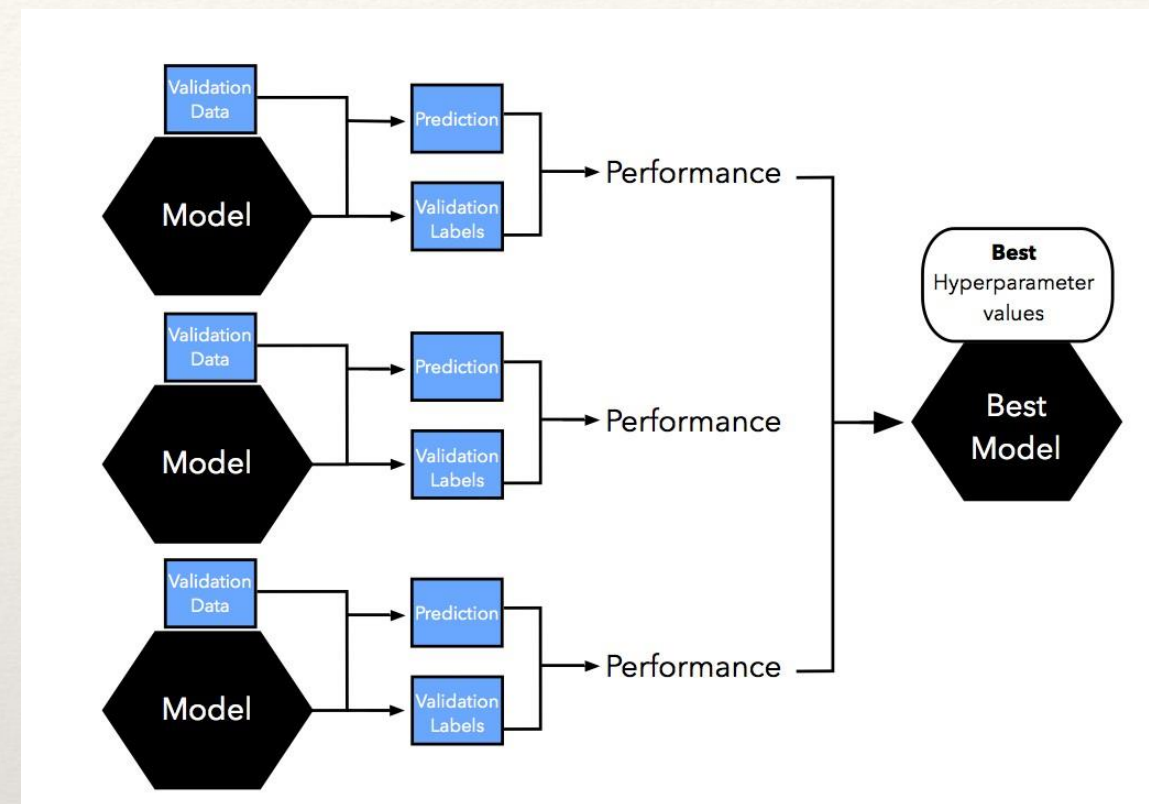
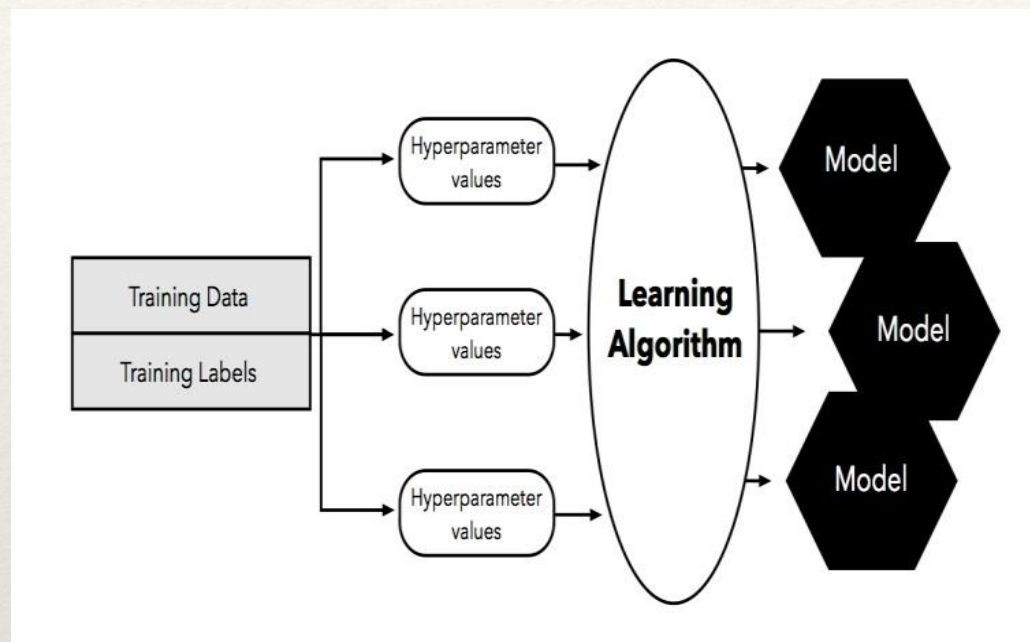
- Split your dataset into 3 disjoint sets: Training, Validation, Test
- If a lot of data are available then you can try 50:25:25 otherwise 60:20:20.



Holdout Method – Hyperparameter tuning

- Identify which parameters need to be optimized
 - e.g., number of hidden neurons, number of hidden layers etc
- Select a performance measure to evaluate the performance on the validation set
 - Accuracy, Precision, Recall etc
 - Appropriate measure depends on the application, if the test set is imbalanced etc


Holdout Method – Hyperparameter tuning



From: <https://sebastianraschka.com/blog/2016/model-evaluation-selection-part3.html>

- Train your algorithm on the training set multiple times, each time using different values for the parameters you wish to optimise.
- For each trained classifier evaluate the performance on the validation set (using the performance measure you have selected).

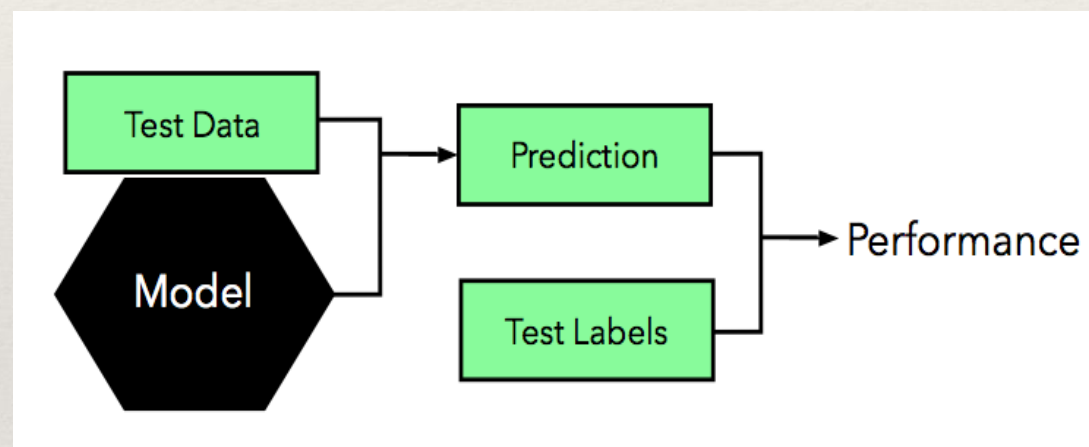
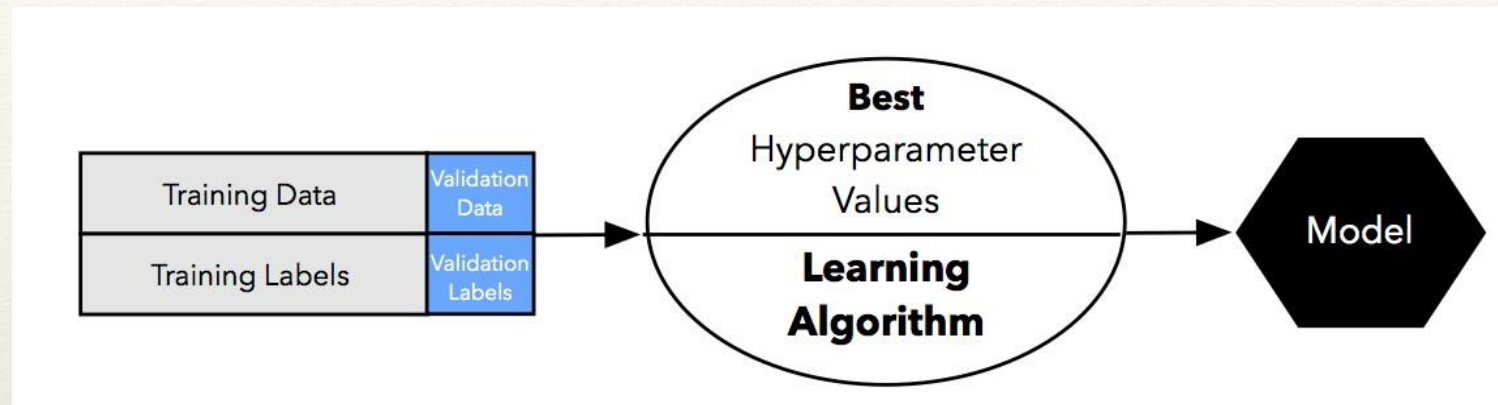
Holdout Method – Hyperparameter tuning



- Keep the classifier that leads to the maximum performance on the validation set (in this example the one trained with 35 hidden neurons).
- This is called parameter optimization/tuning, since you select the set of parameters that have produced the best classifier.

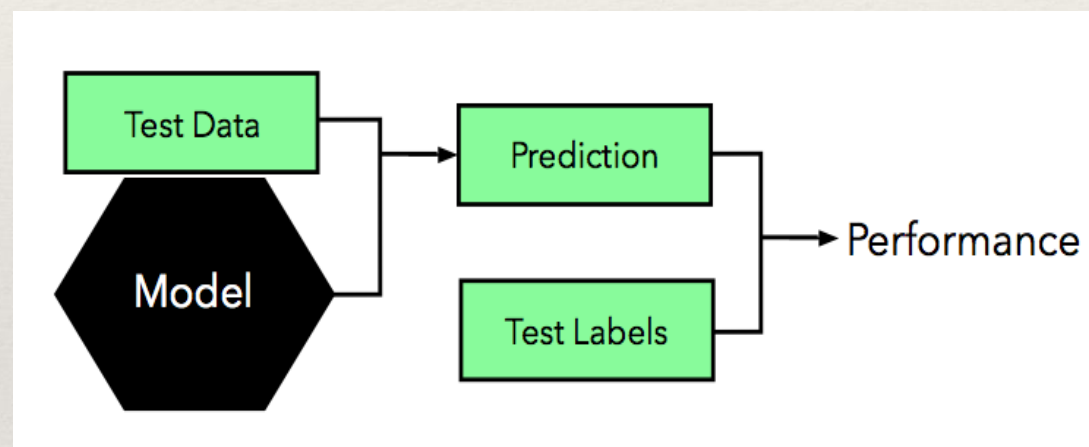
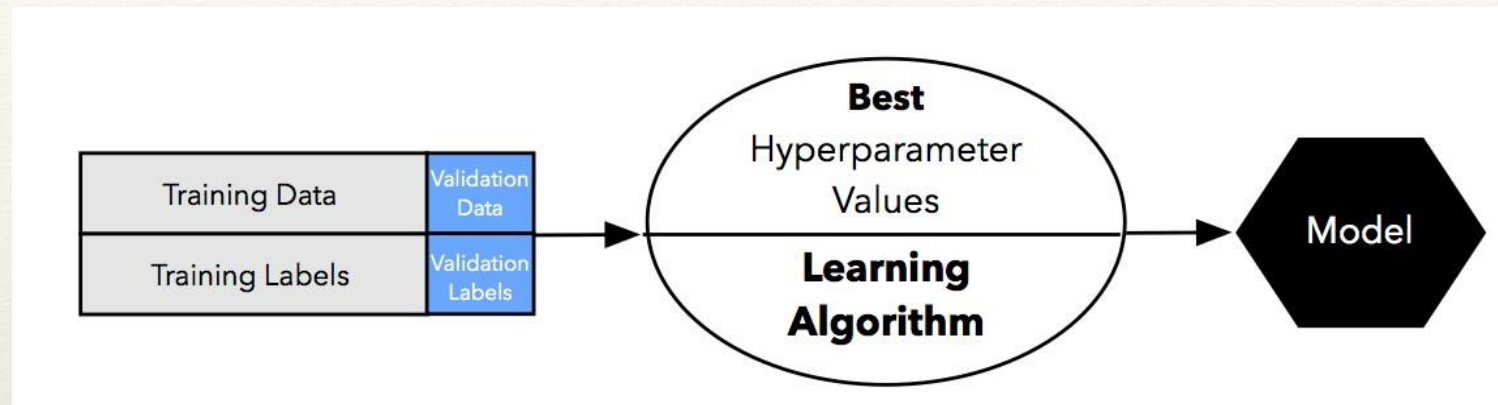
Holdout Method – Hyperparameter tuning

– using parameters in the model



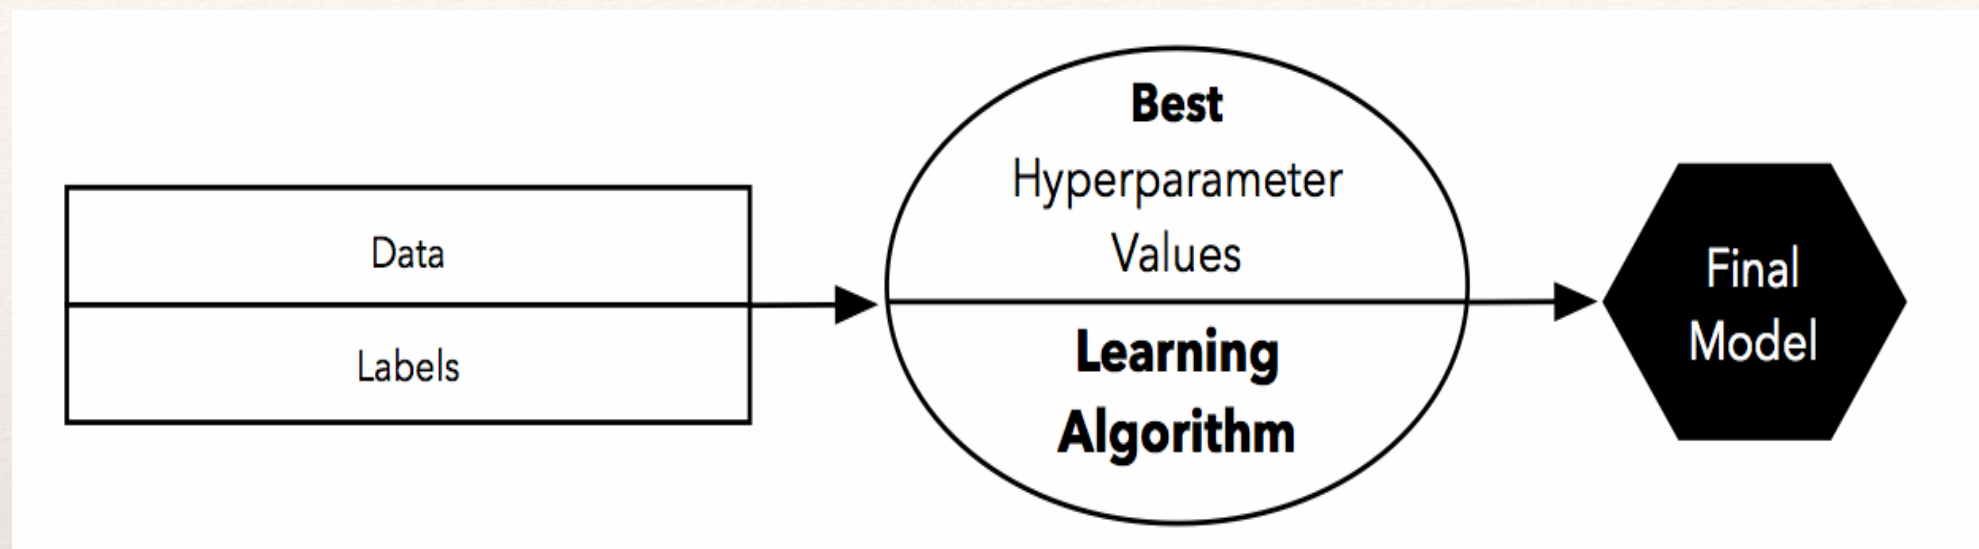
- You can either merge the training and validation sets and train a new classifier using the optimal set of parameters OR you can simply use the best classifier (trained only on the training set).
- Test the performance on the test set.

Holdout Method



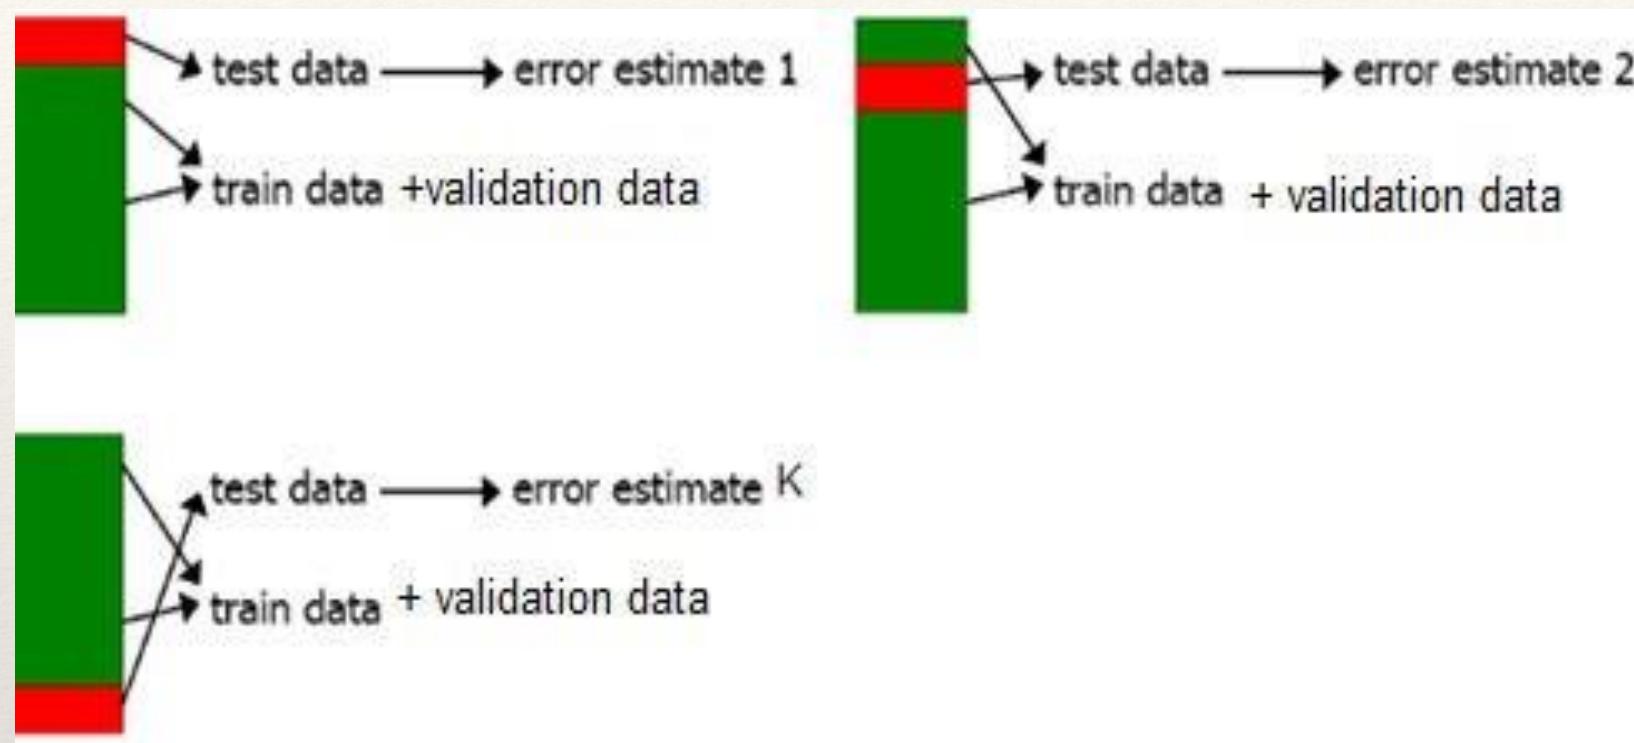
- The test set should **NOT** be used for training or validation. It is used **ONLY** in the end for estimating the performance on unknown examples, i.e. how well your trained classifier generalizes.
- You should assume that you do not know the labels of the test set and only after you have trained your classifier they are given to you.

Holdout Method



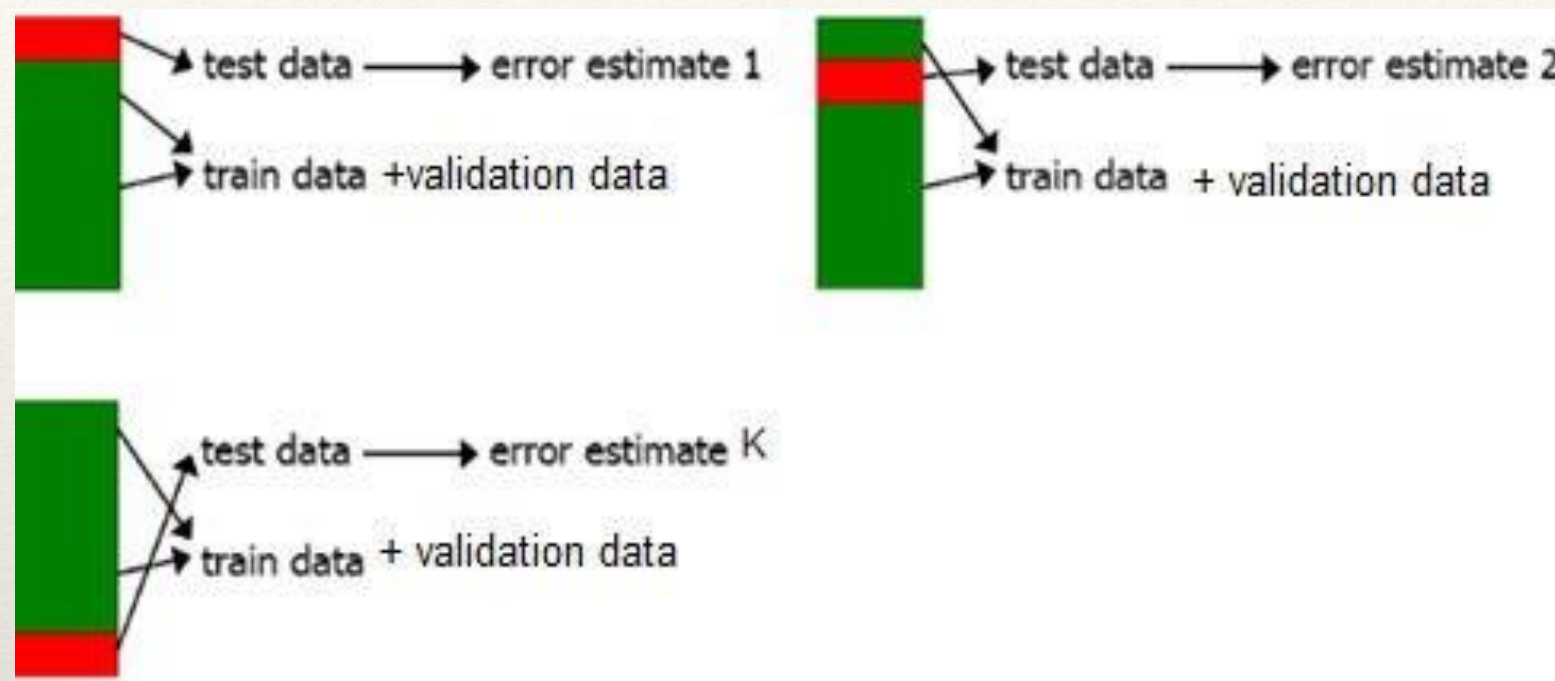
- We need a model which we will use for classifying new examples.
- Either use the one trained on the training set or on training + validation sets OR train a new model on the entire dataset using the optimal set of parameters.

Cross Validation



- When we have a lot of examples then the division into training/validation/test datasets is sufficient.
- When we have a small sample size then a good alternative is cross validation.

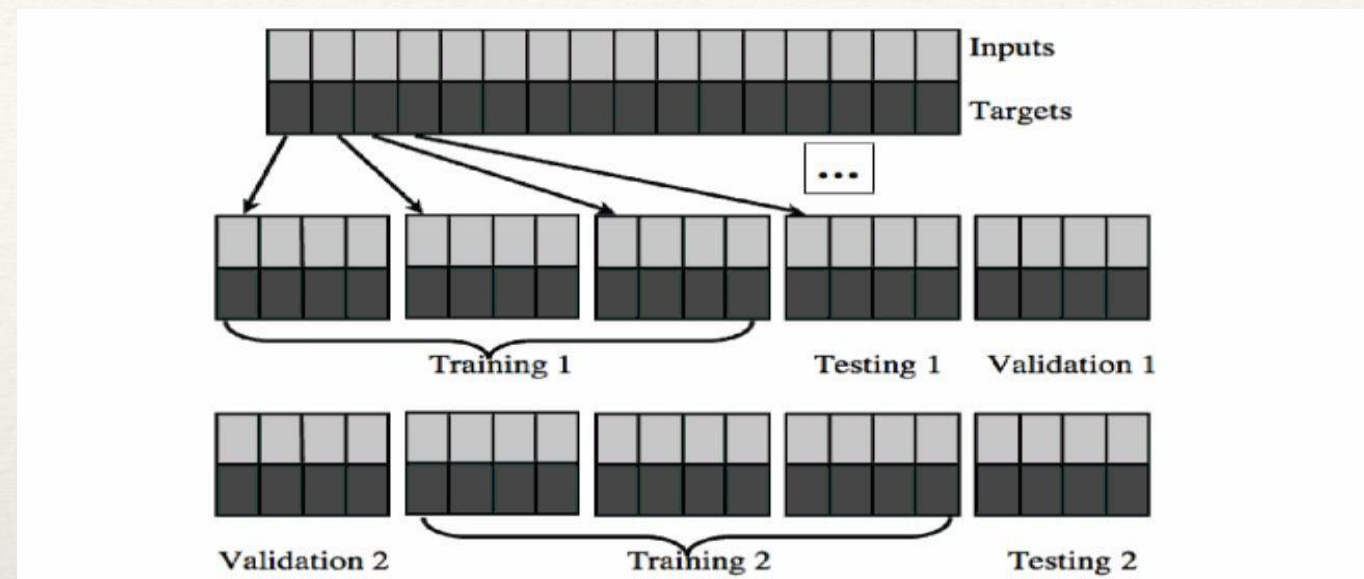
Cross Validation – Test Set Performance Estimation



- Divide dataset into k (usually 10) folds using $k-1$ for training + validation and one for testing
- Test data between different folds should never overlap!
- Training + Validation and test data in the same iteration should never overlap!
- In each iteration the error on the left-out test set is estimated
- Error estimate: average of the k errors

Cross Validation – Test Set Performance Estimation


- The k -folds should be divided into training and validation folds, e.g. k -2 folds for training and 1 for validation.



S. Marsland, Machine learning: An algorithmic perspective

- Train on the training set, optimize parameters on the validation set and test on the test set.
- We can only estimate the test set performance. In other word we evaluate how our implementation (and the way we optimize the parameters) generalizes on unknown test sets.
- We know nothing about the optimal set of parameters. We find a different set of optimal parameters in each fold.

Cross Validation – Parameter Estimation



- We can use cross validation to estimate the optimal set of parameters
- $k-1$ folds for training, 1-fold left out for validation (using the entire dataset)
- For each parameter set run the k fold cross-validation
- Select the parameters that result in the best average performance over all k left out folds

Cross validation – hyperparameter tuning with Grid Search

- ❖ **Grid Search** is a method for adjusting the hyperparameters.
- ❖ With Grid Search, we try all possible combinations of the parameters of interest and find the best ones.
- ❖ Practically – we try best practice values and choose the best combination

Parameter Optimization – Performance Estimation - Summary

- **CASE 1:** A lot of data are available (Holdout Method)
 - 1) Tune parameters on validation set
 - 2) Estimate generalization performance using the test set
 - 3) Train on entire dataset using optimal set of parameters
- **CASE 2:** Data are limited (Cross validation)
 - 1) Run cross validation to estimate the test set performance
 - Training, validation, test folds
 - Optimize parameters in each iteration
 - 2) Run cross validation to estimate optimal parameters
 - Training, Validation folds only (also called Grid Search)
 - 3) Train on entire dataset using optimal set of parameters